

CS 581 – ADVANCED ARTIFICIAL INTELLIGENCE

TOPIC: SEARCH



Mustafa Bilgic



<http://www.cs.iit.edu/~mbilgic>



<https://twitter.com/bilgicm>


MOTIVATION

- AI is powered by search
- Consider multiple possibilities and find the “optimal” one
- Just a few examples
 - Travel from one location to another
 - Find the “optimal” route
 - Game playing
 - Find the “optimal” move
 - Machine learning
 - Find the “optimal” set of parameters
 - Machine translation
 - Find the “optimal” sequence of words
 - Constraint satisfaction
 - Find the “optimal” assignment
 - ...

WE WILL STUDY

- Hill climbing, A*, genetic algorithms, simulated annealing, ...
- Maximum likelihood estimation, Bayesian estimation, expectation maximization, gradient optimization, Lagrange multipliers, policy search, ...
- Alpha-beta search, Monte-Carlo tree search, ...

ABSTRACT REPRESENTATION



Given

- An initial state
- A goal state
- Available actions at each state
- A transition model
- The cost of each action

Find

- A sequence of actions that take you from the initial state to the goal state, where the total cost of the sequence is minimized

TRAVEL PROBLEM

- The world representation
 - In(City)
- Initial state
 - In(Madison, WI)
- Goal state
 - In(Detroit, MI)
- Actions
 - Travel to the neighboring city
- Transition model
 - If you are in city A and travel to city B, the state changes from In(A) to In(B)
- Cost
 - Distance traveled

TRAVELING SALESMAN PROBLEM

- Given
 - A list of N cities
 - Distances between each pair of cities
- Find
 - A minimum-cost travel plan where the agent visits each city exactly once and returns to the city of origin

8 PUZZLE

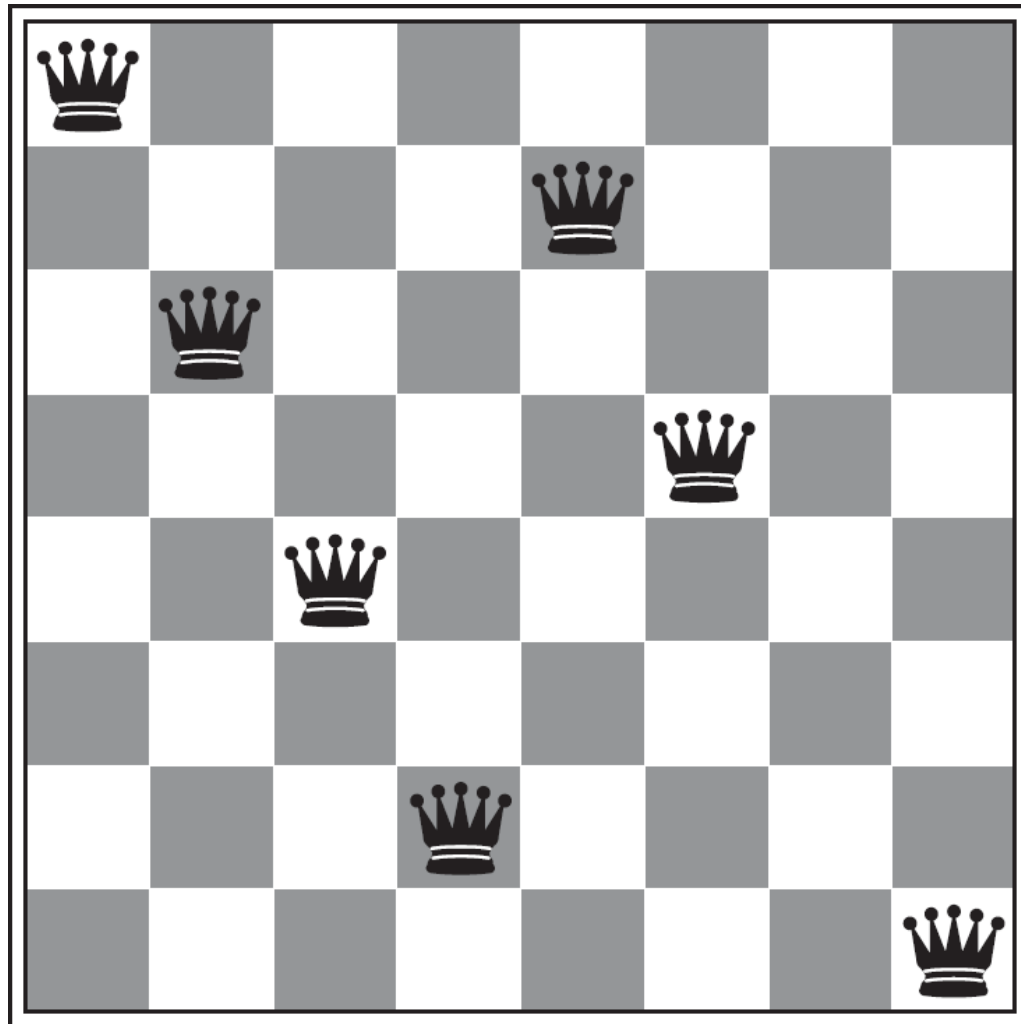
7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

8 QUEENS



OUTLINE

- Uniform cost search, greedy heuristic search, and A* search
 - E.g., travel problem, 3-puzzle problem
- Simulated annealing
 - E.g., traveling salesman problem
- Genetic algorithms
 - E.g., 8-queens problem

EXERCISE

- You are at an initial city D
- You are trying to reach to the city M
- The cities are not fully connected; you can travel from one city to only its neighbor cities
- You don't have access to the full map; you know only the parts you explored or discovered
- Each time you pick a city that you know and ask where you can go from there
- You want to find the shortest path from D to M , and you want to do it as fast as possible
- Notation:
 - $c(n, n')$: the cost of traveling from n to n'
 - $h(n)$: the estimate of the distance from n to the goal
- Three settings:
 - Setting 1: You are given the c function but not h
 - Setting 2: You are given the h function but not c
 - Setting 3: You are given both c and h functions

BEST-FIRST SEARCH PSEUDOCODE

- Given
 - The initial state, goal state, available actions, action costs, and maybe a heuristic function
 - frontier – a priority queue; explored: a set
- Initialize frontier with the initial state
- While frontier is not empty
 - Pop the top node n in the frontier
 - If n is the goal state, return the solution
 - For each child n' of n :
 - If n' is not in explored
 - If n' is not in frontier
 - Add it to frontier
 - Else if the new path to n' is better than the old path, replace the old n' with the new n' in the frontier
- Return failure // frontier is empty; goal was not found

BEST-FIRST SEARCH ALGORITHMS

- Define $f(n) = h(n) + g(n)$
 - $h(n)$: Cost estimate from n to the goal
 - $g(n)$: Cost from the initial state to n
- 1. Uniform-cost search
 - `frontier` is sorted using $g(n)$
- 2. Greedy heuristic search
 - `frontier` is sorted using $h(n)$
- 3. A* search
 - `frontier` is sorted using $f(n)$

A TRAVEL PROBLEM EXAMPLE

- See OneNote

OPTIMAL

- Assume the optimal path cost is p^*
- Uniform cost search is guaranteed to return the optimal path
 - Expands all nodes where $g(n) \leq p^*$
- Greedy heuristic search is not optimal; ignores path costs
- A* is optimal only if h is admissible and consistent
 - Expands all nodes where $g(n) + h(n) \leq p^*$

ADMISSIBLE? CONSISTENT?

○ **Admissible** if

- $h(n)$ never overestimates the optimal cost
- That is $h(n)$ is always optimistic
- E.g., straight line distance between two cities

○ **Consistent** if

- If $h(n) \leq c(n, n') + h(n')$
 - n' is the successor of n
 - Triangle inequality
- If a heuristic is consistent, it is also admissible
(the other way around is not guaranteed)

A* OPTIMALITY – PROOF SKETCH

1. If $h(n)$ is consistent, then $f(n)$ along any path is non-decreasing
2. When n is expanded, the optimal path to it has been found

UCS vs A*

- $A^* = \text{UCS}$ if $h(n) = 0$ for all n
- UCS searches in circles whereas A^* searches in ellipses
- Illustration in OneNote

8 PUZZLE

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

HEURISTICS

- How to design $h(n)$?
 - The answer is “depends on the domain” but some general principles exist
 - For example, relax the problem constraints a bit
- For travel problems
 - Flight distance
- For 8-Puzzle
 - Number of misplaced tiles
 - Manhattan distance

A 3-PUZZLE PROBLEM EXAMPLE

- See OneNote

OTHER SEARCH ALGORITHMS

- Depth-first search
- Breadth-first search
- Iterative deepening search
- Bidirectional search

COMPARING ALGORITHMS

- Time complexity
- Space complexity
- Completeness
 - If there is a solution, can it find it?
- Optimality
 - Is the found solution the optimal one?

COMPARING HEURISTICS

- h_i **dominates** h_j iff
 - $h_i(n) \geq h_j(n) \forall n$
- For the 8-Puzzles problem
 - Manhattan distance dominates # of misplaced tiles
- If you have multiple admissible heuristics where none dominates the other
 - Let $h(n) = \max(h_1(n), h_2(n), \dots, h_m(n))$
 - $h(n)$ is admissible and it dominates all $h_i(n)$
 - Of course, computing a heuristic is not free, and hence we need to consider the computational cost of $h(n)$ compared to each $h_i(n)$

EFFECTIVE BRANCHING FACTOR

- If the total number of nodes generated is N and the solution depth is d , then
 - b^* is the branching factor that a uniform tree of depth d would need to have in order to contain $N+1$ nodes
- $N+1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$
- If A^* finds a solution at depth 4 using 40 nodes, what is b^* ?
 - ≈ 2.182
- A good heuristic function achieves $b^* \approx 1$

EMPIRICAL COMPARISONS

- Figure 3.26 in <http://aima.cs.berkeley.edu/figures.pdf>
- <https://github.com/aimacode/aima-python/blob/master/search4e.ipynb>